AI in 2026: A Tale of Two AIs

By David Cahn

Sequoia Capital | December 2025

Overview

David Cahn -- the Sequoia partner who wrote the landmark '600B Question' piece documenting the gap between AI infrastructure investment and revenue -- returns with a framework for 2026. His central claim: this year will be defined by two contradictory forces operating simultaneously. At the infrastructure and timeline layer, 2026 is the Year of Delays. At the application and adoption layer, it is the Year of Acceleration. Understanding both tracks, and why they coexist, is the essential strategic orientation for anyone operating in AI.

Part I: The Year of Delays

Cahn identifies two distinct categories of delay that will shape the supply side of AI in 2026: physical infrastructure bottlenecks and a reset in AGI timeline expectations.

Data Center Bottlenecks and the TSMC Brake

The hyperscaler buildout -- running at a pace of hundreds of billions in annual capex -- will run headfirst into a supply chain that has not scaled proportionally. The most acute constraint is semiconductor manufacturing. TSMC and ASML occupy monopolistic positions in advanced chip production and cannot be compelled to ramp capacity on demand. As Ben Thompson has named it, the 'TSMC Brake' is structural: TSMC grew revenues 50% since 2022 but grew capex by only 10%, a ratio that reflects considered judgment about sustainable capacity expansion, not timidity.

Data center construction faces a cascade of industrial constraints in its final stages: generators, cooling units, power distribution equipment, skilled construction labor. Dozens of suppliers, each with their own lead times and capacity limits, must coordinate to complete facilities on schedule. Labor shortages in specialized trades needed for large-scale construction have become material. The result is that announced data center capacity does not translate cleanly into deployed capacity -- and in 2026, Cahn expects that gap to be visible.

Practical implication: the infrastructure investment cycle that looked like straightforward demand fulfillment is hitting the limits of physical-world coordination. Announced capex overstates near-term deployable compute.

The AGI Timeline Reset

For several years, 'AGI by 2027' was a casual consensus in parts of Silicon Valley. That consensus has broken down. Cahn identifies a cluster of Dwarkesh Patel podcast interviews --

with Richard Sutton, Andrej Karpathy, and Ilya Sutskever -- as a demarcating line in the discourse. The new center of gravity has shifted: the AGI window is now being discussed in terms of the 2030s at earliest.

This is not a claim that AI progress has slowed. It is a recalibration of what 'AGI' means operationally and how close current systems are to it. The practical effect on enterprise planning: organizations that were in 'wait for AGI' mode need a new posture. The question shifts from 'when will AI be ready?' to 'what can we do with what exists now?'

Part II: The Year of Acceleration

Against the infrastructure delays at the supply layer, Cahn documents accelerating adoption at the application layer. These two tracks are not in contradiction -- they reflect different parts of the stack moving at different speeds.

Enterprise DIY Fatigue Opens the Door for Startups

Large enterprises have spent 2023-2025 attempting to build internal AI capabilities -- standing up data pipelines, fine-tuning models, building RAG infrastructure, assembling AI teams. The experience has been disappointing for many. Internal builds take longer than anticipated, produce narrower capability than promised, and require ongoing maintenance that strains stretched engineering teams.

This 'DIY fatigue' is creating demand for startup solutions. AI startups with vertically integrated, production-ready products are positioned to capture enterprises that tried and struggled to build in-house. The enterprise sales motion shifts from 'here is what you could build with our APIs' to 'here is working software that solves your problem today.'

The $0-to-$1B Club

In 2025, the AI startup ecosystem began tracking a '$0 to $100M ARR' cohort -- companies that reached that revenue threshold faster than any prior generation of software companies. Cahn's prediction for 2026: the relevant benchmark shifts to the $0-to-$1B club.

The structural drivers are the same ones enabling rapid enterprise AI adoption -- proven demand, pre-built infrastructure, compressed sales cycles for urgently needed capabilities -- but at a scale that reflects how large the AI software market has become. The implication for incumbents: competitive threats can emerge and achieve material scale in compressed timeframes that traditional enterprise competitive moats were not designed to defend against.

Synthesis and Implications

Cahn's framework is most useful as a corrective to both extremes of AI discourse. Against AI skeptics who point to infrastructure costs and hype-cycle dynamics as evidence that AI's moment has passed, he documents accelerating real-world adoption. Against AI maximalists who expect the infrastructure buildout to translate immediately into deployable capability, he surfaces the physical-world constraints that introduce meaningful delay.

The strategic orientation this implies: don't wait for infrastructure conditions to improve or AGI timelines to resolve before building on current AI capabilities. The application layer is already running. But don't confuse announced infrastructure investment for near-term deployable capacity -- the gap between capex and compute-in-production is real.

Relationship to Existing Canon

This piece extends Cahn's earlier '$600B Question' -- which identified the gap between infrastructure investment and AI revenue -- with a forward-looking framework for 2026. It complements Sequoia's '2026: This Is AGI' (Grady/Huang) by adding the supply-side and timeline-reset arguments that balance Grady/Huang's bullish adoption thesis. The pairing gives a complete view: adoption is real and accelerating (Grady/Huang); infrastructure timelines and AGI claims need to be stress-tested (Cahn).

The DIY fatigue/startup-solutions argument directly anticipates the organizational transition described in 'From Hierarchy to Intelligence' -- companies will not self-build their way to AI transformation; they will assemble it from specialized, proven solutions.

Source: Sequoia Capital -- https://sequoiacap.com/article/ai-in-2026-the-tale-of-two-ais/
Also published on David Cahn's Substack: https://dcahn.substack.com/p/ai-in-2026-a-tale-of-two-ais
This archival reference was compiled from publicly available summaries and commentary on the original article.