

Policy on the AI Exponential

Dario Amodei, CEO — Anthropic

Published: June 10, 2026

Source: <https://darioamodei.com/post/policy-on-the-ai-exponential>

Executive Summary

Amodei's third major public essay marks a deliberate escalation in Anthropic's policy posture: from advocating transparency legislation to calling for binding, enforceable regulation of frontier AI with real enforcement authority. Timed to coincide with the launch of Claude Fable 5, the essay covers five policy domains and commits Anthropic to \$350 million in new funding to address AI's governance and labor-market impacts. Within two days of publication, the U.S. government used the authority Amodei called for — issuing an export control directive that suspended Fable 5 and Mythos 5 worldwide.

"Frontier AI models, like airplanes, should be required to go through technical testing and auditing, and their release should be blocked or reversed as a threat to public safety if they do not meet high standards of safety."

Why This Essay Is Significant

- From the CEO of Anthropic — a primary voice already in the worldview canon
- The most concrete binding-regulation proposal from any frontier AI lab leader to date
- Directly resolves the worldview's open question: 'How will regulators respond?'
- Explicitly acknowledges that AI job displacement 'may be an intrinsic property of the technology' — a significant shift from earlier optimistic framing
- Government action within 48 hours demonstrates real-world policy influence

Policy Area 1: Safety and Regulation

Mandatory third-party auditing for frontier models trained using more than 10^{25} FLOPs, developed by companies with more than \$500M in AI revenue or \$1B in AI R&D; spending. Audits cover four risk categories:

- Cybersecurity — ability to enable offensive cyberweapons or attacks on critical infrastructure
- Biological weapons — ability to provide meaningful uplift for bioweapon creation
- Loss of control — evidence the model may act autonomously against human intent at scale
- Automated R&D; — ability to autonomously accelerate any of the above three risk categories

If a model fails these audits, government regulators would have authority to block or reverse its deployment — modeled explicitly on FAA authority over aircraft. Amodei argues this framework should be coordinated with allied democracies.

Policy Area 2: Economics and Labor

Amodei breaks from earlier optimism to acknowledge the possibility of large, permanent labor displacement, stating: 'this may be an intrinsic property of the technology and the way it broadly replicates human cognition.' His three-tier policy menu:

- Measurement: Real-time tracking of AI-driven displacement by sector, skill level, geography
- Pro-employment incentives: Wage insurance for workers in lower-paying transition roles; retention tax credits for employers who retrain rather than replace
- Long-term income support: If displacement proves large and permanent, UBI or universal basic services — 'we should not be ideologically opposed to this'

Anthropic committed \$200M to an Economic Futures Research Fund and \$150M to a national fellowship program for early-career Americans displaced by AI.

Policy Area 3: Scientific Acceleration

In downstream fields like biomedicine, the risk is not AI moving too fast — it is regulators moving too slowly. The FDA and EMA should define now what it takes to accept AI-based methods in approval pipelines, before the 7-8 year drug development pipeline becomes the bottleneck. This is 'the one area where we most want speed over caution.'

Policy Area 4: Civil Liberties

- Ban fully autonomous weapons in domestic law enforcement — no AI-controlled lethal force without human authorization
- Close the data broker loophole: prohibit government agencies from purchasing data they could not legally collect directly
- Guarantee AI parity: anyone facing adverse government action must have access to AI assistance at parity with the government's capabilities

Policy Area 5: Geopolitics

Amodei proposes a coalition of democracies to coordinate AI governance. Coalition commitments would include:

- Coordinate AI governance frameworks across member democracies
- Share advanced chips and semiconductor manufacturing equipment among trusted members
- Deny critical supply-chain inputs to adversarial states
- Share AI benefits with developing nations to prevent authoritarian tech dependency
- Reject AI-powered domestic repression as a condition of coalition membership
- Coordinate macroeconomic stabilization as AI drives productivity shocks

Progression from Earlier Amodei Essays

Essay	Date	Core Posture
-------	------	--------------

Machines of Loving Grace	Sept 2024	Optimistic — describes AI's enormous upside potential
The Adolescence of Technology	Jan 2026	Warning — AI will 'test us as a species'; soft governance
Policy on the AI Exponential	June 2026	Enforcement — binding rules, compute thresholds, \$350M commitment

Implications for the AI Worldview

- REGULATION: The worldview's open question 'how will regulators respond?' now has a concrete answer from the frontier: compute-threshold-triggered mandatory audits with deployment blocking authority.
- WORKFORCE: The bear case on jobs has moved from fringe to primary-source acknowledgment. Permanent displacement is now a stated possibility from the CEO of a leading lab.
- GEOPOLITICS: The democratic coalition framing adds substantive structure to the US-China AI race narrative in the worldview.
- BANKING: The civil liberties provisions — data broker restrictions and AI parity in adverse government actions — directly touch how banks interact with regulators and use customer data in underwriting and compliance.
- SAFETY THRESHOLD: The 10^{25} FLOPs threshold establishes a meaningful line relevant for evaluating which AI vendors fall under binding safety requirements.