

The Open-Weights Letter

Jensen Huang, CEO — NVIDIA (with a 25-signatory industry coalition)

Published: July 24, 2026 | Open letter, amplified via Huang's first-ever X post

Source: Open letter signed by 25 companies including NVIDIA, Microsoft, Meta, and Palantir; reporting via Fortune, Axios, TechSpot, BigGo Finance (July 22-24, 2026)

This document is an archival summary reconstructed from the open letter's public quotations and contemporaneous reporting. Huang used this letter to make his first-ever personal post on X, which is itself part of the story.

Executive Summary

Days after Moonshot AI released Kimi K3 — a 2.8-trillion-parameter open-weight model rivaling flagship OpenAI and Anthropic products on some benchmarks — the White House and Treasury began questioning its technical origins and weighing restrictions on Chinese open-weight models over IP and surveillance concerns. In response, Huang broke his own silence on X for the first time ever to back an open letter, signed by 25 companies and organizations (NVIDIA, Microsoft, Meta, and Palantir among them), arguing that restricting open-weight models — including Chinese ones like DeepSeek and Kimi K3 — would repeat a mistake the US narrowly avoided with software in the 1980s. Notably, OpenAI, Anthropic, and Google did not sign.

“AI will transform every industry, power every company, and be built by every country. Open models strengthen safety and cybersecurity, accelerate innovation and diffusion, and enable sovereignty.”

Why This Letter Is Significant

- NVIDIA has no prior standalone presence in the worldview canon despite being arguably the most economically load-bearing company in the entire AI stack — this is Huang's first public policy intervention of essay-length substance, and he chose his literal first personal social post to make it.
- Directly updates the worldview's “Open vs. closed AI” debate row, which currently cites only Zuckerberg (2024 pro-open, 2025 reversal) and LeCun. Huang is a distinct, non-model-lab voice arguing the open side from an infrastructure/hardware vantage rather than a lab vantage.
- The signatory list is the story as much as the content: OpenAI, Anthropic, and Google — the three labs most exposed to competition from cheap, capable open-weight substitutes — all declined to sign, revealing a clean commercial fault line between chip/infrastructure incumbents (who benefit from more inference regardless of model origin) and closed-frontier labs (who benefit from restricting substitutes).
- Substantively engages the worldview's “no one-model-fits-all” belief and vendor lock-in guidance: Huang is explicitly arguing that open-weight models — including foreign ones — should be part of the usable set rather than restricted, which if it prevails as policy would expand rather than narrow the model choice set available to adopters.
- Sits inside a live, high-stakes policy fight: the letter is a direct attempt to shape a pending Washington decision on restricting Kimi K3 and similar models, not commentary after the fact.

The Argument, in Detail

- Historical analogy: draws a direct line to the 1980s, when restrictive export-control proposals for software were considered and ultimately mostly avoided, allowing the US software industry to compound its lead globally. The letter argues restricting open-weight AI models now would repeat the mistake that was avoided then.
- Three claimed benefits of open models: they strengthen safety and cybersecurity (more eyes on weights), accelerate innovation and diffusion (lower barriers to building on top), and enable sovereignty (countries and companies can run inference on infrastructure they control rather than depending on a closed API).
- Huang's Axios interview added specificity beyond the letter itself: he named DeepSeek and Kimi K3 directly, argued the US should not "panic" over China's AI progress, and warned that restricting Chinese open models on national-security grounds could slow American innovation rather than protect it.
- Washington's counter-position: the White House and Treasury have questioned Kimi K3's technical origins, alleging possible distillation-related IP infringement and improper use of export-restricted chips, and are weighing restrictions partly on surveillance-risk and domestic-industry-protection grounds.

Implications for the AI Worldview

- OPEN VS. CLOSED DEBATE: add Huang/NVIDIA as a distinct bull-case voice for openness, motivated by infrastructure economics rather than model-lab philosophy — a useful nuance the current debate row (framed purely as a Zuckerberg/LeCun dispute) does not capture.
- VENDOR STRATEGY: reinforces the worldview's existing "no one-model-fits-all / avoid deep lock-in" guidance — if open-weight models (domestic or foreign) remain available and improve at Kimi K3's pace, the practical menu of viable models for community banks gets wider and cheaper, not narrower, strengthening the case against over-committing to any single closed vendor.
- REGULATORY RISK: a live Washington decision on restricting specific open-weight models is a genuine near-term uncertainty the worldview should flag under "How fast does the technology itself change the calculus?" — a restriction would remove options from the model-choice landscape faster than the worldview's April 2026 baseline anticipated.
- INDUSTRY FAULT LINES: the sharp split between infrastructure players (open) and closed frontier labs (silent/opposed) is a new, concrete data point on how AI-industry incentives are diverging as the market matures — worth tracking as a leading indicator of where competitive and regulatory pressure will land next.