

Pacing the Frontier

1,178 employees across OpenAI, Anthropic, Google DeepMind, and Meta

Published: July 28–29, 2026 | Open letter at pacingthefrontier.com

This document is an archival summary reconstructed from the letter's public quotations and contemporaneous reporting (Fortune, Zvi Mowshowitz, Peter Wildeford, TechTimes, TheNextWeb, Forbes, 80,000 Hours). Direct retrieval of the letter's full text was unavailable at time of archiving; the core request sentence below is quoted consistently and verbatim across independent coverage.

Executive Summary

The letter's central request, quoted directly across coverage:

"We request that the U.S. government support an international effort to develop the technical and governance tools needed to deliberately pace the frontier of automated AI development."

The ask is deliberately narrow: signatories are explicit that they are not calling for a pause or slowdown now — only for building, in advance, the capacity to pace development later if capability growth outruns society's ability to manage it. Coverage ties the letter directly to Anthropic's June 2026 internal research showing AI systems already meaningfully accelerating AI R&D inside labs (recursive self-improvement); the argument is that the tools to pace the exponential need to exist before the moment they're needed. This directly extends the worldview's already-canonized belief from Amodei's "Policy on the AI Exponential" that frontier model access is becoming policy-mediated. Within hours of publication, both OpenAI and Anthropic endorsed the statement as companies, not just as employers of individual signatories — reported as a first: two competing frontier labs jointly backing the same governance ask.

Signatories

- Reported breakdown: roughly 533 Anthropic, 330 OpenAI, 191 Google, and 62 Meta employees, plus others.
- Notable named signatories per coverage: Dario Amodei (CEO, Anthropic — already a canon voice), Jared Kaplan, Chris Olah, and Jack Clark (Anthropic co-founders); Jakub Pachocki (Chief Scientist, OpenAI) and Mark Chen (Chief Research Officer, OpenAI); Shengjia Zhao and Dawn Song (Meta AI); Anca Dragan, Shane Legg, and Jasjeet Sekhon (Google DeepMind); and Ilya Sutskever (CEO, Safe Superintelligence Inc. — already tracked in the worldview's AGI-timeline debate row).

Critical Reception

- Christian Catalini (Forbes, "Don't Pace The Frontier. Look Inside The Trojan Horse", July 29, 2026) argues international pacing agreements are structurally unenforceable: treaties can't be verified, compute accumulates in the shadows, and export-control pressure has already helped trigger the domestic breakthroughs it was meant to prevent in China. His alternative is investment

in verification and audit infrastructure — the ability to actually inspect what's happening inside AI development — rather than reliance on unenforceable pledges to slow down.

- Lance Eliot (Forbes) notes the letter's “support” ask is vague and non-prescriptive about the actual mechanism, leaving open how any pacing infrastructure would function in practice.
- A recurring critique across commentators: without independent verification, statements like this are “cheap talk.” Some read this letter — like Hassabis's standards-body proposal — as a competitive move dressed as a safety measure: raising the barrier behind incumbent labs at the expense of future entrants and open-weight competitors. That reading sits in direct, unresolved tension with the open-weights coalition already in canon (Huang's letter), which several of the same corporate actors are adjacent to on the opposite side of the openness question.

Implications for the AI Worldview

- KEY DEBATE — Model access: market or policy: adds employee-level, cross-lab evidence to a row currently built on CEO-level statements and government export actions. That over a thousand technical staff across four rival labs signed the same statement is a different kind of evidence than another individual essay — it suggests the policy-mediation belief reaches beyond leadership messaging into the rank and file.
- Extends the already-canonized belief that frontier model access is becoming policy-mediated rather than only commercially negotiated — this is the most direct evidence yet that the shift is broadly held inside the labs themselves, not just asserted by their CEOs.
- Sits in genuine, unresolved tension with the Huang open-weights coalition already in canon: the worldview's Key Debates table should probably note this contradiction explicitly rather than treat the two rows as independent.
- OPEN QUESTION — “How far does agent autonomy actually go inside a regulated institution, and who signs when nobody reviewed?”: pacing infrastructure of the kind this letter asks for is precisely the category of governance tool that would need to exist before Sequoia's Stage 3 (“dark factory”) autonomy could be considered safe in a regulated setting like banking.