

A US-Led International AI Forum

Sam Altman, CEO — OpenAI

Published: July 1, 2026 | Financial Times (op-ed)

Source: Financial Times, "Sam Altman: We need a US-led international forum for AI" (July 1-2, 2026)

This document is an archival summary reconstructed from the Financial Times op-ed and contemporaneous reporting (Fortune, Forbes, SiliconANGLE, Gizmodo, Newsweek). The original op-ed sits behind the FT paywall; content here is drawn from extensive direct quotation in secondary coverage published within 48 hours of release.

Executive Summary

Published four days after the G7 summit in France — where OpenAI, Anthropic, and Google DeepMind executives met with world leaders to discuss common standards for frontier models — Altman's op-ed is OpenAI's first concrete governance proposal at this level of specificity. It calls for a "US-led international forum" combining government representatives and independent technical experts to set safety standards, assess capabilities and risk, and gate access to frontier AI technology on compliance. The op-ed also lands three weeks after OpenAI's recapitalization into a public benefit corporation and amid an active OpenAI push for a US government equity stake in the company — giving the proposal a self-interested as well as civic reading.

"AI will reshape the material conditions of human life on a scale that no technology has accomplished since the harnessing of electricity, and perhaps beyond even that."

"Democratic institutions must not cede their responsibilities to AI labs. The labs develop the technology, but citizens and their elected representatives must make the rules."

Why This Op-Ed Is Significant

- From the CEO of OpenAI — a primary voice already anchoring the worldview's Altman track ("The Gentle Singularity," "Three Observations") — but this is his first essay-length governance framework, not a capability or product forecast.
- Creates a direct, comparable counterpart to Amodei's "Policy on the AI Exponential" (already in canon): two frontier-lab CEOs have now published competing regulatory blueprints within three weeks of each other — Amodei's is domestic and binding (compute-threshold audits, deployment-blocking authority); Altman's is international and access-gated (participate-and-comply to earn frontier access).
- Directly engages the worldview's open question, "How will regulators respond?" — by proposing the answer come from a new supranational body rather than existing frameworks (SR 11-7, EU AI Act) alone.
- China's government publicly rejected the framing (calling Altman's democracy-vs-authoritarianism framing "groundless"), confirming the proposal is already a live geopolitical fault line, not a hypothetical.
- Enforcement is the acknowledged weak point: unlike aircraft or enrichment plants, frontier models train inside opaque data centers, making external verification of compliance genuinely unresolved — coverage notes this as the proposal's central unanswered question.

The Proposal, in Detail

- Composition: government representatives plus independent technical experts, not labs alone — explicit reversal of the "let the labs self-govern" posture OpenAI has sometimes been associated with.
- Function: establish accepted safety standards, provide impartial analysis of capabilities and risk, and make frontier technology available to nations and companies that participate and follow the rules.
- Precedent models cited: aviation safety regulation, global financial standards bodies, and the International Atomic Energy Agency (IAEA).
- Framing: AI development as a contest between democratic and authoritarian governance models, with the US positioned to lead the democratic coalition — echoing (but distinct from) Amodei's "coalition of democracies" geopolitics proposal.

Implications for the AI Worldview

- REGULATION DEBATE: Adds a second, distinct data point to the worldview's "AI regulation approach" debate row — both leading US lab CEOs (Amodei, Altman) are now on record favoring more structure than the deregulation camp (Trump EOs, Andreessen), but they disagree on binding domestic enforcement vs. an access-gated international body. That intra-industry split is itself worth tracking.
- OPEN QUESTION #3 ("How will regulators respond?"): the worldview can now note that frontier labs are trying to shape the answer proactively rather than waiting for regulators to act — a meaningfully different posture than the document's April 2026 baseline assumed.
- GEOPOLITICS: reinforces the worldview's US-China AI race narrative, but China's on-record rejection of the framing is a useful corrective against treating "coalition of democracies" language as settled or uncontested internationally.
- BANKING: largely indirect — no banking-specific provisions (unlike Amodei's civil-liberties/data-broker proposals). Relevant mainly as context for how frontier-model access and export controls could evolve, which would affect which vendors community banks can rely on long-term.