

Built to Benefit Everyone: Our Plan

Authors: Sam Altman (CEO) and Jakub Pachocki (Chief Scientist) — OpenAI

Published: June 8, 2026

Source: <https://openai.com/index/built-to-benefit-everyone-our-plan/>

Companion development: Sam Altman, Financial Times op-ed proposing a US-led international AI governance forum, published on or about July 1, 2026, following the June 17, 2026 G7 summit.

This document was compiled from publicly available reporting and analysis of the original essay and its companion op-ed for inclusion in the AI World View canon. Both original pieces (openai.com and the Financial Times) returned access-denied errors during compilation; this summary is triangulated across multiple independent outlets that quote the primary text directly (TechRadar, Forbes, Axios, Gizmodo, SiliconANGLE, theprint, AI Weekly) and should be read as a faithful secondary account, not a verbatim copy. The original essay is available at the source URL above.

Overview

Published June 8, 2026, 'Built to Benefit Everyone: Our Plan' is OpenAI's most complete public statement of strategy since Altman's 2025 'Gentle Singularity' essay, and the company's clearest attempt to answer the question its critics keep asking: benefit everyone, or entrench the company that got there first? Altman and chief scientist Jakub Pachocki frame OpenAI as entering a 'third phase' and use a deliberate historical analogy — rural electrification in the 1920s — to argue that AI access is fundamentally an infrastructure-equity problem: a technology this consequential cannot be allowed to concentrate in a small number of institutions.

Three Near-Term Goals

- **Automated AI researcher.** Build an AI system that can accelerate and increasingly automate the research process itself, while remaining steerable, accountable, and connected to people. OpenAI states an internal expectation that by March 2028 a significant fraction of its own research could be done by AI systems working in tandem with human researchers.
- **Broad-based economic growth.** Accelerate scientific progress, productivity, and economic growth while working to ensure the resulting gains are widely shared, not concentrated among early adopters and capital owners.
- **Personal AGI for every person on Earth.** Give each individual access to a personal AGI, empowering them to use it however they choose — the essay's most direct restatement of OpenAI's founding mission.

The Power-Distribution Argument

The essay's structural claim: 'a good AI future cannot be one where a small number of institutions control most of the capability.' Altman and Pachocki argue from historical pattern — concentrated power creates fragility, while widely shared power makes societies more resilient, adaptable, and free — and use it to justify wide distribution of both access and, implicitly, governance authority.

The essay pairs this with an explicit guardrail against full automation: 'entirely automating everything is not the future we want,' which the authors describe as both unfulfilling and dangerous. 'Powerful systems must remain safe, aligned with human intent, and subject to human control' — people should

continue to set goals, exercise judgment, and decide what is worth pursuing, even as execution is automated.

Secondary reporting (Forbes, EdTech Innovation Hub) frames this around five stated principles guiding how OpenAI builds and deploys AI: democratization, empowerment, universal prosperity, resilience, and adaptability.

Companion Proposal: A US-Led International AI Forum

Roughly three weeks later, in a Financial Times op-ed timed after the June 17, 2026 G7 summit, Altman extended the governance argument into a concrete geopolitical proposal: a 'US-led international forum that establishes accepted standards, provides expert and impartial analysis of capabilities and risks, and makes the technology available to nations and companies that participate and follow the rules.' He cited global aviation safety standards, international financial standards, and the IAEA as the closest historical precedents.

Altman argues explicitly that democratic institutions, not AI labs, must retain the rule-making authority: 'the most important decisions about how AI is used should be made through democratic processes — not by a small number of companies in San Francisco.'

The same reporting window (Axios, Forbes, Gizmodo, Time) describes OpenAI in parallel talks with the Trump administration — Commerce Secretary Howard Lutnick and Treasury Secretary Scott Bessent — about a potential ~5% US government equity stake in OpenAI (roughly \$42.6B on an \$852B valuation), contingent on similar arrangements from competing labs. This sits in tension with the essay's own anti-concentration argument and is worth reading skeptically: a voluntary, US-chaired standards forum paired with a bespoke equity relationship between one lab and one government is a distribution-of-power claim that cuts against itself.

Contrast With Amodei's 'Policy on the AI Exponential' (Already in Canon)

This piece is best read alongside Dario Amodei's June 10, 2026 essay already in this canon. The two documents stake out visibly different governance philosophies from the two most prominent frontier labs, published two days apart:

- **Anthropic (Amodei):** binding, compute-threshold-triggered regulation (10^{25} FLOPs) with real enforcement authority — explicitly modeled on FAA authority to ground aircraft, coordinated with allied democracies, backed by \$350M of Anthropic's own funding for the transition.
- **OpenAI (Altman/Pachocki):** a voluntary, US-led international forum that sets standards and grants access to compliant participants — modeled on aviation safety bodies and the IAEA, but without a stated enforcement or deployment-blocking mechanism, and paired with a proposed direct equity stake for the US government specifically.

Both essays invoke aviation-safety analogies to argue for oversight; they diverge sharply on whether that oversight is binding and multilateral (Anthropic) or voluntary and US-chaired (OpenAI). That divergence is itself a first-order fact for anyone tracking how AI governance actually develops over the next 1–2 years, and is a cleaner encapsulation of the 'AI regulation approach' debate than either

essay alone.

Implications for the AI Worldview

- **Open Question – ‘How will regulators respond?’** Now has two competing frontier-lab answers on record, not one. The worldview's existing note (from the Amodei addition) that this question ‘has a concrete answer from the frontier’ should be revised to reflect that the frontier is not unified: binding multilateral audits (Anthropic) vs. voluntary US-led standards-setting plus a national equity stake (OpenAI).
- **Debates to track table:** worth adding a row — ‘AGI governance model’: bull case (Altman/Pachocki — voluntary US-led forum, wide distribution of access) vs. bear case (Amodei — binding compute-threshold enforcement; also: critics noting the OpenAI proposal's anti-concentration rhetoric sits awkwardly next to a proposed US-government equity stake in OpenAI specifically).
- **AGI timeline debate:** the March 2028 internal target for AI systems doing ‘a significant fraction’ of OpenAI's own research is a new, dated, falsifiable claim — useful as a concrete marker alongside Aschenbrenner (2027), Amodei (2–3 years), and the METR task-doubling curve already in the canon (Sequoia, ~2028 crossover).
- **Belief #17 (AI-native challengers) and #62 (world model):** the ‘personal AGI for everyone’ goal is a direct bet that individual-level AI access, not institutional context alone, will determine who benefits — a mild complication to this worldview's emphasis on proprietary institutional context as the scarce resource. Worth holding both ideas in tension rather than resolving them.
- **Not directly banking-specific:** unlike the BCG/McKinsey banking reports scanned in this same cycle, this pair of documents operates at the macro AI-governance layer, alongside Altman's Gentle Singularity and Amodei's two essays already in the canon. Its relevance to Alloy Labs is indirect — it shapes the regulatory and competitive environment banks will eventually operate under — rather than offering a banking-specific data point.

Recommendation: add to canon. Do not treat as requiring an immediate rewrite of the worldview document itself — its main effect is to sharpen the existing open question on regulatory response and add a concrete new entry to the debates-to-track table, both of which can be folded into the next full worldview revision rather than an ad hoc edit.