# AI Ascent 2026: This Is AGI

Pat Grady, Sonya Huang, Konstantine Buhler — Sequoia Capital

Published: May 2026

Source: https://sequoiacap.com/article/ai-ascent-2026/

## Executive Summary

At its fourth annual AI Ascent conference, Sequoia's three lead AI partners declared that from a commercial standpoint, AGI has arrived. The definition is functional and deliberately non-philosophical: a system dispatched to perform a job, capable of recovering from failure, that persists until the job is complete. The essay introduces several new frameworks: the METR curve, the diffusion gap, the MAD founder playbook, and a three-stage model of agent autonomy culminating in the 'dark factory.'

> *"When you can dispatch an agent to execute a task and it can recover from failure and persist until the job is done, that is functional AGI. By that definition, 2026 is the year it arrived."*

## The Three Inflection Points

Grady identifies three discontinuous capability moments:

- November 2022 — ChatGPT reveals pre-training scale to the mainstream
- Late 2024 — OpenAI o1 surfaces a second scaling law: inference-time compute. Reasoning models prove that compute at inference (not just training) scales performance.
- 2026 — Claude Code / Fable generation demonstrates long-horizon agents: systems that pursue goals across hours and days, not just tokens. The third scaling law is persistence.

## 2026: The Year of Agents

Three ingredients have converged to make capable agents viable at scale:

### 1. The Brain (Models)

Leaps in reasoning capability and context length allow agents to plan and self-correct. Models now hold complex multi-step tasks in context and reason through them without losing coherence across long interaction sequences.

### 2. The Limbs (Tools)

MCP (Model Context Protocol), terminal access, web search, and computer use enable agents to execute actions in software environments, call APIs, and manipulate files — not just generate text but operate in the real world.

### 3. Resilience (Harness)

Reinforcement learning enables agents to recover from failure. Prior agents collapsed at the first unexpected state; RL-trained agents treat failure as information and adapt. This is the ingredient that transforms a powerful assistant into an autonomous worker.

## The METR Curve: AGI Timeline by Data

The METR (Model Evaluation for Trustworthiness and Reliability) score measures the maximum duration of an autonomous task an AI agent can reliably complete. The score doubles every 7 months. At this rate:

- Early 2026: agents reliably complete tasks measured in hours
- ~2028: agents reach human-equivalent task duration — the crossover point
- At crossover: an agent is functionally equivalent to a full-time employee, but never sleeps, never takes days off, and can be infinitely replicated at marginal cost

The METR curve is significant because it grounds AGI timeline estimates in measured capability doubling rates rather than speculation. It implies workforce-equivalent AI is approximately two years away, not a decade.

## Three Stages of Agent Autonomy

- Stage 1 — Supervised: A human directs one agent or coordinates a small team. Human reviews every significant action before it executes.
- Stage 2 — Background / Async: Agents operate without continuous human supervision and report results asynchronously. Human sets objectives and reviews outputs.
- Stage 3 — Dark Factory: Human review is removed from the production loop entirely. Agents spin up sub-agents, review each other's work, and complete tasks end-to-end.

Dark factory operations are already appearing at cybersecurity companies doing automated penetration testing. The conditions under which Stage 3 is safe vs. reckless are 'not settled' — this is the frontier's open governance problem.

## The $10 Trillion Services Opportunity

The software industry has never been able to address the $10 trillion global professional services market. Capable agents change this equation — they can perform complex, judgment-intensive work previously exclusive to human professionals:

- Medicine: genome analysis, intervention recommendations, prescriptions, clinical trials
- Law: contract negotiation, litigation support, dispute settlement
- Science: solving Erdos-class mathematics, proposing novel superconductors
- Consumer: inbox management, calendar coordination, financial oversight, tax filing
- Finance: analysis, modeling, compliance monitoring, report generation

The economic activity once captured by professional service firms is being absorbed into the application layer. 'Services is the new software' — a thesis Sequoia first introduced in 2025 and now considers confirmed by agent capabilities.

## The Diffusion Gap

A widening gap exists between the rate at which new AI capabilities are created in labs and the rate at which they reach enterprise deployment. Foundation models advance daily; the average Fortune 500 absorbs change in months or years. Every day the gap widens, the application-layer opportunity grows for founders and operators who can bridge it.

For community banks specifically, the diffusion gap is structurally wide: legacy cores, regulatory overhead, and risk aversion create unusually slow adoption. This is simultaneously the greatest risk (competitors close the gap faster) and the greatest opportunity (the application-layer value for banking is enormous and largely uncaptured).

## MAD: The Founder Playbook for the Agent Era

- M — Moats from the customer back: Build competitive advantages in customer relationships, proprietary data, and workflow integration — not model superiority (commoditizing rapidly)
- A — Affordance design: Design for what AI can uniquely do, not what humans did before. Don't automate old workflows; reimagine them around agent capabilities.
- D — Diffusion gap exploitation: Target industries where the gap between frontier capability and current deployment is largest — that is where application-layer value accrues.

## Implications for the AI Worldview

- AGI FRAMING: The functional/commercial AGI definition resolves the philosophical debate. Sequoia's definition is now measurable via the METR curve — not a matter of opinion.
- WORKFORCE: The METR curve provides a specific, data-grounded timeline for agent-human capability parity (~2028). The worldview should track this metric explicitly.
- BANKING: The diffusion gap maps directly to the worldview's banking analysis — the widest diffusion gaps create the largest application-layer value opportunities.
- SERVICES MODEL: Quantifies the 'selling work not seats' disruption at $10T. The SaaS disruption thesis in the worldview is now anchored to a specific market size.
- DARK FACTORY: The three-stage agent autonomy model gives banks a practical framework for their AI adoption journey from supervised pilots to autonomous operations.
- METR AS MEASUREMENT: Banks should track METR scores for deployed agents as a leading indicator of when autonomous operations become viable in each function.